Boolean Gröbner Basis Reductions on Finite Field Datapath Circuits using the Unate Cube Set Algebra

Utkarsh Gupta, Priyank Kalla, *Senior Member, IEEE*, Vikas Rao

Abstract—Recent developments in formal verification of arithmetic datapaths make efficient use of symbolic computer algebra algorithms. The circuit is modeled as an ideal in polynomial rings, and Gröbner basis (GB) reductions are performed over these polynomials to derive a canonical representation. As they model logic gates of the circuit, the ideals comprise largely of Boolean (or pseudo-Boolean) polynomials. This paper considers a logic synthesis analogue of GB reductions over Boolean polynomials, by interpreting symbolic algebra as the unate cube set algebra over characteristic sets. By representing Boolean polynomials as characteristic sets using Zero-suppressed BDDs (ZDDs), implicit algorithms are efficiently designed for GB-reduction for datapath circuits. We show that the imposition of circuit-topology based monomial orders exposes a special structure on the ZDD representation of the polynomials. The subexpressions employed in the GB-reduction are readily visible as subgraphs on the ZDDs, which are directly used to compose the result. Our division algorithms effectively cancel multiple monomials implicitly in one-step, simplify the search for divisors, and avoid intermediate size explosion. Experiments performed over various finite field arithmetic architectures demonstrate the efficiency of our algorithms and implementations; our approach is orders of magnitude faster as compared to conventional methods.

I. Introduction

Automated formal verification and equivalence checking of arithmetic datapath circuits is challenging. Conventional verification techniques, such as those based on binary decision diagrams (BDDs) [1], And-Invert-Graph (AIG) based reductions with SAT or SMT-solvers [2], etc., are infeasible in verifying complex datapath designs. Such designs often implement algebraic computations over bit-vector operands, therefore finite integer rings [3] [4] or finite fields [5] [6] are considered appropriate models to devise decision procedures for verification. For this reason, the verification community has explored the use of algebraic geometry and symbolic algebra algorithms for verification.

In such a setting, the logic gates of the circuit are modeled by way of a set of multivariate polynomials $F = \{f_1, \ldots, f_s\}$ in rings $R[x_1, \ldots, x_n]$. Usually, the coefficients $R = \mathbb{Z}, \mathbb{Z}_{2^k}$, or $\mathbb{F}_{2^k}$, depending on whether the integer, finite integer ring (mod 2^k), or respectively the finite field (of 2^k elements) model is employed for verification. This set of polynomials F generates an ideal, and for verification it is required to compute a *Gröbner basis (GB)* [7] of this ideal. Reducing the primary output polynomials of the circuit modulo this GB results in a unique canonical polynomial expression, and it can be used for equivalence checking.

The GB problem exhibits high computational complexity. Indeed, computing a GB for large circuits is practically infeasible. Managing this complexity ought to be a major goal of any approach.

1) State-of-the-art & Limitations: Recent approaches [3] [5] have discovered that particularly for circuit verification problems, the expensive GB computation can be avoided altogether. For arbitrary combinational [3] [5] and sequential circuits [8], a specialized term order $>$ can be derived by analyzing the topology of the given circuit. This term order is derived by performing a reverse topological traversal of the circuit, and in this manuscript we refer to it as the *Reverse Topological Term Order* (RTTO). Imposition of RTTO $>$ on the polynomial ring *renders the set of polynomials of the circuit itself a GB*. Subsequently, the verification problems can be solved solely by way of GB-reduction (using multi-variate polynomial division), without any need to explicitly compute a GB. It has now become standard practice to make use of RTTO-style term orders to solve various formal verification problems on digital circuits (see for example [3], [5], [4], [9], [10], [11]), where the early techniques of [3] [5] have been extended and improved to verify integer and floating point arithmetic circuits [9], [11], [12], [13].

A common theme among all these relevant works is that *by virtue of RTTO, they move the complexity of verification from one of computing a GB to that of GB-reduction by way of multivariate polynomial division.* Moreover, since the Gröbner basis is derived from the logic gates of the circuit, *it comprises Boolean polynomials.* Boolean polynomials (formally defined in Section III) consist of terms that have coefficients from $\mathbb{F}_2 = \{0, 1\}$, and monomials that are a product of variables where the degree of each variable is also restricted to the set $\{0, 1\}$. The aforementioned approaches will benefit greatly by a *dedicated, domain-specific implementation of GB-reduction w.r.t. Boolean polynomials, carried out on the given circuit under RTTO $>$*. So far, the above techniques [3], [6], [5], [11], [9], [14] use a general-purpose polynomial division approach, together with an explicit representation, for this GB-reduction. Moreover, the overall concept of polynomial division is still utilized in its rudimentary form, involving iterative cancellation of monomials "1-step at a time" on explicit data-structures. Despite recent efforts, such GB-reductions can lead to *a worst-case size explosion problem*, which needs to be addressed.

2) Objective & Rationale: This paper addresses the problem of deriving canonical representations for datapath circuits as Boolean polynomials. These canonical Boolean polynomials are used for equivalence checking of datapath designs. The canonical representation requires a Gröbner basis reduction modulo a set of Boolean polynomials. To make this GB-

reduction on circuits more efficient, this paper describes new techniques, algorithms and implementations, specifically targeted for circuit verification under RTTO $>$. We make use of the *implicit* characteristic set representation for storing and manipulating Boolean polynomials using *Zero-Suppressed BDDs (ZDDs)* [15]. By analyzing the structure of ZDDs for polynomial representation under RTTO $>$, we show how this GB-reduction can be efficiently implemented using algorithms that specifically manipulate the ZDD graph, by interpreting Boolean polynomial manipulation as the *algebra of unate cube sets.*

The algebraic objects used to model the polynomial ideals derived from digital circuits are rings of Boolean polynomials. When Boolean functions are represented in $\mathbb{F}_2$ (AND/XOR expressions), and that too as a canonical Gröbner basis, the representation tends to explode. Polynomial representations employed in computer algebra tools, such as the *dense-distributive data-structure* of the SINGULAR computer algebra tool [16], are inefficient for this purpose. Since addition and multiplication (mod 2) are equivalent to XOR and AND operations, respectively, GB-reduction can be viewed as a polynomial analog of a specialized *AND/XOR Boolean decomposition* problem. Moreover, the monomials of a Boolean polynomial are a product of literals in positive polarity, which can be viewed as *unate cubes* in logic synthesis. Clearly, implicit Boolean set representations such as decision diagrams could be employed for Gröbner basis reductions over Boolean polynomials.

3) Technical Contributions: We first describe when and how the GB-reduction encounters a term-explosion (exponential blow-up) under RTTO $>$, which cannot be easily overcome by explicit representations. We show that ZDDs can avoid this exponential blow-up – thereby justifying their use. Subsequently, *we show that RTTO $>$ imposes a special structure on ZDDs that allows to implement reduction techniques that implicitly cancel multiple monomials in every step of polynomial division.* Moreover, due to RTTO $>$, the subexpressions that are required for polynomial division are also readily available as subgraphs in the ZDDs. Our algorithm exploits this special structure, thus improving GB-reduction in both space and time.

Using an implementation integrated with the CUDD [17] package, we perform extensive experiments on datapath circuits for deriving the canonical representation of the functions implemented by them. The benchmark designs include various cryptography primitives, such as finite field multipliers, elliptic curve point addition circuits, and also sequential finite field circuits. Experiments conducted on these benchmarks show *orders of magnitude improvement* using our implementation of GB-reduction, as compared against contemporary methods. In fact, for these benchmarks, our bit-level (Boolean) approach is much faster than the word-level approaches (e.g. [6]).

4) Paper Organization: The following section reviews relevant previous work on Gröbner basis based verification of datapath circuits, and the literature on Boolean Gröbner basis and applications. Section III describes the mathematical background on Gröbner basis reductions, the RTTO based term order $>$, and how these concepts are applied to datapath veri-

fication. Section IV motivates how and why the Boolean GB-reduction can be viewed as the unate cube set algebra. Section V describes the theory, algorithms and implementations for GB-reduction for Boolean polynomials using ZDDs. Section VI describes the experiments conducted for verification using our approach, and Sec VII concludes the paper.

II. RELATED PREVIOUS WORK

In the past decade, computer algebra and algebraic geometry based datapath verification has received a lot of attention. The work of [18] used a GB-reduction approach to derive canonical representations of (word-level) RTL datapath descriptions over finite rings $\mathbb{Z}_{2^k}$. Using the same finite ring model, [3] addressed data correctness properties of arithmetic bit-level implementations using a Gröbner basis formulation. This paper showed how an efficient term order $>$ can be derived from the circuit to simplify the computation.

In [5], the authors addressed formal verification of finite field arithmetic circuits using the Strong Nullstellensatz formulation over $\mathbb{F}_{2^k}$. Using a set (F) of polynomials to describe the logic circuit, along with a set of vanishing polynomials (F_0) over the field $\mathbb{F}_{2^k}$, the verification problem was formulated as a (radical) ideal membership test, requiring a Gröbner basis. Drawing inspirations from [3], the authors in [5] also exploited the same concept of deriving a specialized term order $>$ to simplify the ideal membership test. In particular, it was shown that $>$ could be derived by performing a *reverse topological traversal* on the circuit. Imposition of this term order $>$ rendered the set of polynomials $F \cup F_0$ itself a Gröbner basis, and verification was then performed simply by a GB-reduction. This GB-reduction was subsequently formulated as Gaussian elimination on a coefficient matrix [5], [6] in the style of the F_4 algorithm; called F_4-style reduction in the sequel.

Formulations of a similar flavor (GB-reduction under the specialized term order $>$ derived from the circuit) were used and integrated with SMT-solvers [4] for verification using a pseudo-Boolean model (akin to $\mathbb{Z}_{2^k}[X] \pmod{X^2 - X}$). More recently, these concepts have been applied to verify integer arithmetic [9], [11], [13], and also floating point circuits [12]. The authors in [14] show that the reduction process can be *parallelized* by performing reduction for each output bit independently.

While polynomial division algorithms form the core computation in the above techniques, the overall concept of polynomial division is still utilized in its rudimentary form, involving iterative cancellation of monomials "1-step at a time" on explicit data-structures.

A. Boolean Gröbner Basis

The symbolic algebra community has studied properties of Boolean Gröbner bases [19] [20]. Boolean GB formulations have also been used for SAT solving [21], e.g. to derive proof refutation, and for model checking [22] [23]. From among these, the work of PolyBori [20] comes closest to ours, and is a source of inspiration for this work. PolyBori proposed the use of ZDDs to compute Gröbner bases for Boolean polynomials. PolyBori is a *generic* Boolean GB computational

engine that caters to many permissible term orders. Its division algorithm is also generally based on the conventional concept of canceling one monomial in every step of reduction. In contrast, our algorithms are tailored for GB-reduction under the RTTO $>$. The efficiency of our approach stems from the observation that the RTTO $>$ imposes a special structure on the ZDDs, which allows for multiple monomials to be canceled in one division-step, along with simplifying the search for divisors. Experiments show that our approach is an order of magnitude faster than PolyBori.

III. BACKGROUND: GRÖBNER BASIS REDUCTION AND CANONICAL REPRESENTATIONS

Let $\mathbb{B} = \{0,1\}$ denote the Boolean domain, $\mathbb{F}_2$ the finite field of 2 elements ($\mathbb{B} \equiv \mathbb{F}_2$), and $R = \mathbb{F}_2[x_1,\ldots,x_n]$ the polynomial ring over variables $x_1,\ldots,x_n$ with coefficients in $\mathbb{F}_2$. Operations in $\mathbb{F}_2$ are performed (mod 2), so $-1 = +1$ in $\mathbb{F}_2$. We will use $+,\cdot$ to denote addition and multiplication in R, and $\neg, \vee, \wedge$ and $\oplus$ to denote Boolean negation, OR, AND and XOR operations, respectively.

A polynomial $f \in R$ is written as a finite sum of terms $f = c_1 X_1 + c_2 X_2 + \cdots + c_t X_t$. Here $c_1,\ldots,c_t$ are coefficients and $X_1,\ldots,X_t$ are monomials, i.e. power products of the type $x_1^{e_1} \cdot x_2^{e_2} \cdots x_n^{e_n}$, $e_i \in \mathbb{Z}_{\geq 0}$. To systematically manipulate the polynomials, a monomial order $>$ (also called a term order) is imposed on the ring. This order $>$ is a total order and a well order on all the monomials of R such that multiplication by a monomial preserves the order[1]. All polynomials in R are represented using $>$. Subject to $>$, when a polynomial is written as $f = c_1 X_1 + c_2 X_2 + \cdots + c_t X_t$ such that $X_1 > X_2 > \cdots > X_t$, we call $lt(f) = c_1 X_1$, $lm(f) = X_1$, $lc(f) = c_1$, the *leading term*, *leading monomial* and *leading coefficient* of f, respectively, with $lt(f) = lc(f) \cdot lm(f)$. We also denote tail$(f)$ $= f - lt(f) = c_2 X_2 + \cdots + c_t X_t$. In this work, we are mostly concerned with terms ordered lexicographically (*lex*).

Definition III.1 (Boolean Polynomial). *Let $f = c_1 X_1 + \cdots + c_t X_t$ be a polynomial in $\mathbb{F}_2[x_1,\ldots,x_n]$ such that the coefficients $c_i \in \{0,1\}$, and monomials $X = x_1^{e_1} \cdot x_2^{e_2} \cdots x_n^{e_n}, e_i \in \{0,1\}$. Then f is called a* **Boolean polynomial**. *For Boolean polynomials $lt(f) = lm(f)$.*

A gate-level circuit can be modeled with Boolean polynomials, where every Boolean logic gate operator is mapped from $\mathbb{B}$ to a polynomial function over $\mathbb{F}_2$:

$$
\begin{aligned}
z = \neg a &\mapsto z + a + 1 \quad (\text{mod } 2) \\
z = a \wedge b &\mapsto z + a \cdot b \quad (\text{mod } 2) \\
z = a \vee b &\mapsto z + a + b + a \cdot b \quad (\text{mod } 2) \\
z = a \oplus b &\mapsto z + a + b \quad (\text{mod } 2)
\end{aligned}
\tag{1}
$$

Polynomial reduction via division: Let f, g be polynomials. If $lm(f)$ is divisible by $lm(g)$, then we say that f is *reducible to r modulo g*, denoted $f \xrightarrow{g} r$, where $r = f - \frac{lt(f)}{lt(g)} \cdot g$. This operation forms the core operation of polynomial division algorithms and it has the effect of canceling the leading term

---

of f. Similarly, f can be *reduced w.r.t. a set of polynomials* $F = \{f_1,\ldots,f_s\}$ to obtain a remainder r. This reduction is denoted as $f \xrightarrow{F}_+ r$, and the remainder r has the property that no term in r is divisible (i.e. cannot be canceled) by the leading term of any polynomial f_i in F. Algorithm 1 (Alg. 1.5.1 from [7]) depicts the procedure to perform this classical reduction that cancels one monomial in every iteration of the while-loop.

Algorithm 1 Multivariate Reduction of f by $F = \{f_1,\ldots,f_s\}$

1: **procedure** *multi_variate_reduce*$(f, \{f_1,\ldots,f_s\}, f_i \neq 0)$
2:     $u_i \leftarrow 0$; $r \leftarrow 0$; $h \leftarrow f$
3:     **while** $h \neq 0$ **do**
4:         **if** $\exists i$ s.t. $lm(f_i) \mid lm(h)$ **then**
5:             choose i least s.t. $lm(f_i) \mid lm(h)$
6:             $u_i = u_i + \frac{lt(h)}{lt(f_i)}$
7:             $h = h - \frac{lt(h)}{lt(f_i)} f_i$
8:         **else**
9:             $r = r + lt(h)$
10:            $h = h - lt(h)$
11:     **return** $(\{u_1,\ldots,u_s\}, r)$

Polynomial ideals: Given a set of polynomials $F = \{f_1,\ldots,f_s\}$ from the ring $R = \mathbb{F}_2[x_1,\ldots,x_n]$, the ideal generated by F is $J = \langle F \rangle \subseteq R$:

$$
J = \langle f_1,\ldots,f_s \rangle = \{h_1 \cdot f_1 + \cdots + h_s \cdot f_s : h_1,\ldots,h_s \in R\}, \tag{2}
$$

where the polynomials $f_1,\ldots,f_s$ are called the generators (or basis) of the ideal.

For a binary variable x_i, we have $x_i^2 = x_i$. Therefore, the polynomial $x_i^2 - x_i$ vanishes over $\mathbb{F}_2$, and we call it a vanishing polynomial. While manipulating Boolean polynomials it is essential to ensure this indempotency by reducing the polynomials modulo (the ideal of) all vanishing polynomials $\langle x_1^2 - x_1,\ldots,x_n^2 - x_n \rangle$. Therefore, we essentially operate on the quotient ring of $\mathbb{F}_2[x_1,\ldots,x_n]$ modulo the ideal of vanishing polynomials, i.e. over $\mathbb{F}_2[x_1,\ldots,x_n]/\langle x_1^2 - x_1,\ldots,x_n^2 - x_n \rangle$. Then Boolean polynomials are exactly the canonical representatives of residue classes in the aforementioned quotient ring. A significant benefit of using a Boolean data-structure such as ZDDs is that the reduction $x_i^2 = x_i$ is performed implicitly by the data-structure. For this reason, in the sequel we *omit explicit mention* of reduction modulo the vanishing polynomials and it should be assumed that the reduction $x_i^2 = x_i$ is always performed.

Gröbner basis of ideals: An ideal J may have many different sets of generators, i.e. it is possible to have $J = \langle f_1,\ldots,f_s \rangle = \langle h_1,\ldots,h_r \rangle = \cdots = \langle g_1,\ldots,g_t \rangle$. A Gröbner basis G of ideal J is one such set of polynomials $G = GB(J) = \{g_1,\ldots,g_t\}$ with many important properties that allow to solve many polynomial decision questions.

Definition III.2 (Gröbner basis [7]). *For a monomial ordering $>$, a set of non-zero polynomials $G = \{g_1, g_2,\ldots,g_t\}$ contained in an ideal J, is called a Gröbner basis of J iff $\forall f \in J$, $f \neq 0$, there exists $g_i \in \{g_1,\ldots,g_t\}$ such that $lm(g_i)$ divides $lm(f)$; i.e., $G = GB(J) \Leftrightarrow \forall f \in J : f \neq 0, \exists g_i \in G : lm(g_i) \mid lm(f)$.*

---

[1]Lexicographic (*lex*) and degree-lexicographic (*deglex*) are examples of such permissible monomial orders.

Gröbner basis G of an ideal $J = \langle f_1, \ldots, f_s \rangle$ is computed using the Buchberger's algorithm [24], reproduced in Alg. 2.

Algorithm 2 Buchberger's Algorithm

Inputs: $F = \{f_1, \ldots, f_s\}$
Outputs: $G = \{g_1, \ldots, g_t\}$
1: $G := F$;
2: **repeat**
3:    $G' := G$
4:    **for** each pair $\{f_i, f_j\}, i \neq j$ in G' **do**
5:       $Spoly(f_i, f_j) \xrightarrow{G'}_+ h$
6:       **if** $h \neq 0$ **then**
7:          $G := G \cup \{h\}$
8: **until** $G = G'$

In the algorithm,

$$Spoly(f_i, f_j) = \frac{L}{lt(f_i)} \cdot f_i - \frac{L}{lt(f_j)} \cdot f_j \qquad (3)$$

where $L = LCM(lt(f_i), lt(f_j))$. An important property of a Gröbner basis G is that reduction of a polynomial f modulo G is essentially unique.

Theorem III.1 (Gröbner Basis Reduction, Thm. 1.9.1 in [7]). *Let $G = \{g_1, \ldots, g_t\}$ be a Gröbner basis of ideal J, and let f be another polynomial. The reduction of f modulo G, denoted $f \xrightarrow{G}_+ r$, is called the **Gröbner basis reduction (GBR)** of f. The remainder r so obtained by GBR of f is a canonical expression modulo G.*

In other words, for any polynomial f, if $f \xrightarrow{G}_+ r_1$ and $f \xrightarrow{G}_+ r_2$, then $r_1 = r_2 = r$. The remainder r is sometimes also called the *normal form* of f w.r.t. G. The canonicity of r (modulo G) can be exploited for equivalence checking of digital circuits.

Proposition III.1. *Given a circuit C, we can represent all the gates using (Boolean) polynomials $F = \{f_1, \ldots, f_s\}$ in $\mathbb{F}_2[x_1, \ldots, x_n]$ by means of Eqn. (1), and generate ideal $J = \langle F \rangle$. Let $z_i, \ i = 0, \ldots, k-1, (z_i \in \{x_1, \ldots, x_n\})$ denote one-bit of the k-bit primary output variables of the circuit. Compute a Gröbner basis $G = GB(J) = \{g_1, \ldots, g_t\}$ for the polynomials of the circuit, and perform the GBR $z_i \xrightarrow{G}_+ r_i$ for all $0 \leq i < k$. Then all r_i's are a canonical representation and can be used for formal verification/equivalence checking.*

To derive the canonical representation r_i, it is required to compute a Gröbner basis G of the ideal generated by the polynomials of the circuit. Buchberger's algorithm for computation of a Gröbner basis exhibits high complexity. Over finite fields $\mathbb{F}_q$ of q elements ($q = 2$ in our case), the complexity is bounded by $q^{(O(n))}$, where n is the number of variables [25]. This complexity still needs to be overcome. The work of [5] showed that for formal verification of combinational circuits, the expensive GB computation can be avoided altogether, due to the following results.

Lemma III.1 (Product Criterion [26]). *For two polynomials f_i, f_j in any polynomial ring R, if the equality $lm(f_i) \cdot lm(f_j) =$*

$LCM(lm(f_i), lm(f_j))$ *holds, i.e. if $lm(f_i)$ and $lm(f_j)$ are relatively prime, then $Spoly(f_i, f_j) \xrightarrow{G}_+ 0$.*

Using this criterion we can say that when the leading terms of *all polynomials in the basis* $F = \{f_1, \ldots, f_s\}$ are relatively prime, then all $Spoly(f_i, f_j) \xrightarrow{G}_+ 0$. As no new polynomials are generated in Buchberger's algorithm, F is already a Gröbner basis ($F = GB(J)$). For a combinational circuit C, a specialized term order $>$ can always be derived by analyzing the circuit topology which ensures such a property [3] [5]:

Proposition III.2. *(From [5]) Let C be an arbitrary combinational circuit. Let $\{x_1, \ldots, x_n\}$ denote the set of all variables (signals) in C. Starting from the primary outputs, perform a reverse topological traversal of the circuit and order the variables such that $x_i > x_j$ if x_i appears earlier in the reverse topological order. Impose a lex term order $>$ to represent each gate as a polynomial f_i, s.t. $f_i = x_i + tail(f_i)$. Then the set of all polynomials $\{f_1, \ldots, f_s\}$ forms a Gröbner basis G, as $lt(f_i) = x_i$ and $lt(f_j) = x_j$ for $i \neq j$ are relatively prime. This term order $>$ is called the **Reverse Topological Term Order (RTTO)**.*

Imposition of RTTO on the polynomials of the circuit has the effect of making every gate output variable x_i a leading term of f_i. Since every gate output is unique, $lm(f_i) = x_i, lm(f_j) = x_j$ become relatively prime. As a result, the set F is already a GB ($G = F$), the explosive GB computation is avoided, and verification is performed solely by the canonical GB-reduction: $z_i \xrightarrow{G}_+ r_i$. Note that as $f_i = x_i + tail(f_i)$, RTTO ensures that every variable x_j that appears in $tail(f_i)$ satisfies $x_i > x_j$. Moreover, the remainder r_i comprises only primary inputs of the circuit. These properties will be exploited in our algorithms.


Fig. 1: A 2-bit modulo Multiplier circuit.

$$f_1 : c_0 + a_0 \cdot b_0, \ lm = c_0; \quad f_2 : c_1 + a_0 \cdot b_1, \ lm = c_1$$
$$f_3 : c_2 + a_1 \cdot b_0, \ lm = c_2; \quad f_4 : c_3 + a_1 \cdot b_1, \ lm = c_3$$
$$f_5 : r_0 + c_1 + c_2, \ lm = r_0; \quad f_6 : z_0 + c_0 + c_3, \ lm = z_0$$
$$f_7 : z_1 + r_0 + c_3, \ lm = z_1$$

Fig. 2: Polynomials of the circuit under RTTO constitute a GB.

Example III.1. Demonstration of the approach: *Consider the circuit given in Fig. 1. Impose RTTO on the circuit. The primary outputs z_0, z_1 are both at level-0, variables r_0, c_0, c_3 are at level-1, c_1, c_2 are at level-2, and the primary inputs*

a_0, a_1, b_0, b_1 *are at level-3. Order the variables* $\{z_0 > z_1\} > \{r_0 > c_0 > c_3\} > \{c_1 > c_2\} > \{a_0 > a_1 > b_0 > b_1\}$. *Using this variable order, we impose a lex term order on the monomials. Then all the polynomials extracted from the circuit have relatively prime leading terms, as shown in Fig. 2, and* $G = F = \{f_1, \ldots, f_7\}$ *forms a GB.*

Then the GBRs $z_1 \xrightarrow{G}_+ a_0 \cdot b_0 + a_1 \cdot b_1$ *and* $z_0 \xrightarrow{G}_+ a_0 \cdot b_1 + a_1 \cdot b_0 + a_1 \cdot b_1$ *are canonical expressions (Boolean polynomials) of the output bits and can be used for equivalence checking.*

IV. Unate Cube Sets & Boolean Polynomials

A Boolean *variable* represents a dimension of the Boolean space $\mathbb{B}^n$, and a *literal* is an instance of a variable x_i or its complement $\neg x_i$. A *cube* is a product of literals which denotes a point or a set of points in the Boolean space. A *cube set* consists of a collection of cubes, each of which is a combination of literals. *Unate cube sets* allow for the use of only positive literals, not negative/complemented literals.

When cube sets are used to represent Boolean functions, they are usually *binate* cube sets containing negative literals. In binate cube sets, literals x_i and $\neg x_i$ represent $x_i = 1$ and $x_i = 0$, respectively; while the absence of a literal implies a *don't care*. In unate cube sets, literal x_i implies $x_i = 1$ whereas its absence implies $x_i = 0$. For example, the cube set $\{a, bc\}$ corresponds to points $(abc) : \{111, 110, 101, 100, 011\}$ in the binate cube set representation, whereas it represents $(abc) : \{100, 011\}$ in the unate cube set representation.

Each monomial of a Boolean polynomial can be viewed as a unate cube – a product of positive literals – and a Boolean polynomial as a unate cube set. Then the GBR $z_i \xrightarrow{G}_+ r_i$ can be interpreted as operations over unate cube sets, as shown below. Let us (re)consider the one-step division for Boolean polynomials: $f \xrightarrow{g} r$. This division can be interpreted as:

$$f \xrightarrow{g} r = f - \frac{lt(f)}{lt(g)} \cdot g \qquad (4)$$

$$= f - \frac{lm(f)}{lm(g)} \cdot g; \quad (\ lt(f) = lm(f)) \qquad (5)$$

$$= f + \frac{lm(f)}{lm(g)} \cdot g; \quad (-1 = +1 \pmod 2) \qquad (6)$$

$$= f \oplus \frac{lm(f)}{lm(g)} \cdot g; \quad (as + \pmod 2 = \oplus) \qquad (7)$$

Notice that $\frac{lm(f)}{lm(g)}$ is a unate product of literals. i.e. a unate cube. The $\oplus$ operation cancels common cubes from f and $\frac{lm(f)}{lm(g)} \cdot g$. The $\cdot$ operation models the modulo 2 product, where the partial products are added using the $\oplus$ operation. Therefore, we will implement the reduce operation $f \xrightarrow{g} r$ for Boolean polynomials as $r = f \oplus \frac{lm(f)}{lm(g)} \cdot g$ using an implicit representation particularly suited for such unate cube operations, i.e. ZDDs.

A. Zero Suppressed Binary Decision Diagrams (ZDDs) for Boolean Polynomials

Binary decision diagrams (BDDs) [1] and their variants have been used as implicit representations for solving many Boolean and pseudo-Boolean optimization problems. Zero-suppressed BDDs [15] are another variant of BDDs that were designed to efficiently manipulate "sets of combinations". A ZDD is obtained from an ordered BDD by: i) eliminating those vertices whose 1-edge (then-edge) is incident on the 0-terminal; and ii) merging isomorphic subgraphs. Subject to the given variable order, a ZDD represents a Boolean function canonically. Just as with BDDs, every node in a ZDD is assigned an index, which corresponds to the variable order imposed on the diagram. For a detailed description of ZDDs and their capabilities for solving logic optimization and sparse combinatorial problems, the reader is referred to [15] and [27].


Fig. 3: ZDD for the polynomial $r_1 = yd + y + d$.

In [27], *Minato* demonstrated that ZDDs are an efficient data-structure for implicit manipulation (algebra) of unate cube sets. Fig. 3 depicts a ZDD for the unate cube set $\{yd, y, d\}$ with the variable order $y > d$. The paths beginning from the root node y and terminating in the 1-terminal node correspond to the cubes of the set. A variable is present in a cube if its 1-edge lies on the path; otherwise it is absent from the cube if its 0-edge lies on the path. Consequently, the ZDD of Fig. 3 can be construed to represent the Boolean polynomial $r_1 = yd + y + d$. *Minato* has shown ([15][27]) how the set union, intersection and difference operations can be implemented recursively on the ZDDs, and they have been implemented using the *ite-operator* in decision diagrams such as the CUDD [17] package. We extend these operations to accommodate the sum and product operations (mod 2), i.e. polynomial algebra in $\mathbb{F}_2[x_1, \ldots, x_n]$, by manipulating sets of combinations using ZDDs.

V. Theory and Algorithms

Let us first consider a scenario where the GBR $z \xrightarrow{G}_+ r$ on circuits under RTTO can result in a size explosion (expression swell) problem using an explicit representation. We also show on the other hand that an implicit set representation (ZDD) overcomes this size explosion.


Fig. 4: A chain of OR gates.

Consider the circuit shown in Fig. 4 consisting of a chain of OR gates. Impose RTTO: i.e. *lex* term order with variable ordering as, $z > y > x > d > c > b > a$. The Boolean polynomials
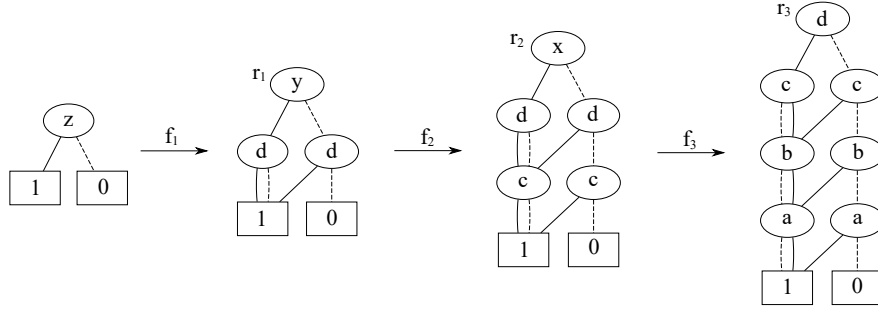
Fig. 5: Reduction of output of the circuit in Fig. 4 by f_1, f_2, f_3.

for the circuit are:

$$f_1 = z + yd + y + d; \qquad (8)$$
$$f_2 = y + xc + x + c; \qquad (9)$$
$$f_3 = x + ba + b + a; \qquad (10)$$

Under RTTO, the set $G = \{f_1, f_2, f_3\}$ forms a GB. To derive a canonical representation of the function, we have to reduce the output $z \xrightarrow{G}_+ r$. A classical symbolic algebra reduction using an explicit representation is carried out as follows:

1) $z \xrightarrow{f_1} yd + y + d$
2) $yd + y + d \xrightarrow{f_2} y + xdc + xd + dc + d \xrightarrow{f_2} xdc + xd + xc + x + dc + d + c$
3) $xdc + xd + xc + x + dc + d + c \xrightarrow{f_3} xd + xc + x + dcba + dcb + dca + dc + d + c \xrightarrow{f_3} xc + x + dcba + dcb + dca + dc + dba + db + da + d + c \xrightarrow{f_3} x + dcba + dcb + dca + dc + dba + db + da + d + cba + cb + ca + c \xrightarrow{f_3} dcba + dcb + dca + dc + dba + db + da + d + cba + cb + ca + c + ba + b + a = r$

Observations: i) Notice that the size of the final remainder corresponds to that of the worst case of a Boolean polynomial: i.e. r contains $2^n - 1 (= 15)$ monomial terms for $n (= 4)$ variables. ii) Classical division algorithms reduce the polynomials 1-step at a time, where only one monomial is canceled in each step. iii) The number of 1-step reductions can increase exponentially as GBR progresses across the circuit.

It is clear that any data-structure that *explicitly* represents each monomial will encounter space and time explosion: this includes the dense-distributive representation of SINGULAR computer algebra tool [16], or the ones used by [9], [11]. The F_4-style polynomial reduction of [5], [6] simulates division on a matrix M representing the problem. However, each column of M corresponds to a monomial generated in the division process, therefore [5], [6] also encounter this size explosion.

The use of ZDDs can help overcome this explosion. Fig. 5 shows the same reduction of z by f_1, f_2, f_3 using ZDDs (exact procedure discussed later). The size of the ZDDs after complete reduction by f_1, f_2, f_3 increases linearly in the number of nodes. Subsequently, the final remainder has $2 \cdot n - 1 (= 7)$ nodes (excluding the terminal 1 and 0 nodes) for $n (= 4)$ variables. Notice that while this controls space explosion, the number of paths (monomials) in the ZDD is still exponential in the number of variables. A classical division algorithm that cancels only one monomial at a time may still

require an exponential number of iterations. We show how to improve upon such a situation.

Problem Setup: Given a circuit C, denote all its *nets* with variables $x_1, \ldots, x_n$. Let the total number of gates in the circuit be s; represent each gate of the circuit with a polynomial f_i in its immediate inputs. Then $F = \{f_1, \ldots, f_s\}$ describes the circuit netlist as a set of Boolean polynomials in $\mathbb{F}_2[x_1, \ldots, x_n]$.

Objective: Impose RTTO on the polynomial ring, so that the set $F = \{f_1, \ldots, f_s\}$ constitutes a Gröbner basis G. For all primary outputs $z_i \in \{x_1, \ldots, x_n\}$, compute $z_i \xrightarrow{G}_+ r_i$ where r_i is the canonical representation of z_i modulo G, and use it for equivalence checking. The representation of the polynomials F of C, and the computation $z_i \xrightarrow{G}_+ r_i$ is to be performed using ZDDs.

A. ZDD representation for the polynomials of the circuit

First, a reverse topological traversal of the circuit is performed to derive the variable order $x_1 > x_2 > \cdots > x_n$ as given in Prop. III.2. The same variable order is imposed on the ZDDs, i.e. x_1 is the variable at the top level in the ZDDs. A ZDD is created for each variable x_i. Using Eqn. (1) the gates of the circuit are modeled as the set of Boolean polynomials $F = \{f_1, \ldots, f_s\}$. ZDDs for these polynomials are constructed using the $+$ and $\cdot$ binary operations for modulo 2 sum and product of two ZDDs. Conceptually, the modulo 2 sum ($\oplus$) operation for two ZDDs f, g can be implemented as $f + g = f_{cs} \cup g_{cs} - f_{cs} \cap g_{cs}$, where f_{cs} and g_{cs} represent the cube sets for the polynomials f and g respectively. For example, let $f = ab + c$ and $g = c + d$ with the corresponding cube sets $f_{cs} = \{ab, c\}$ and $g_{cs} = \{c, d\}$, then $f_{cs} \cup g_{cs} = \{ab, c, d\}$ and $f_{cs} \cap g_{cs} = \{c\}$. The set difference $f_{cs} \cup g_{cs} - f_{cs} \cap g_{cs}$ is the set $\{ab, d\}$ and the corresponding Boolean polynomial is $ab + d$.

Experience has shown that such an implementation with the union operation results in large size of intermediate ZDDs. In order to avoid this intermediate size explosion, we have implemented the $f + g \pmod 2$ operation along similar lines as presented in [20]. The algorithm for this operation is shown in Algorithm 3.

Example V.1. *To demonstrate $f + g, f = ab + c, g = c + d$, let the variable ordering be $a > b > c > d$, i.e. index values for these variables are $0, 1, 2, 3$, respectively. The condition $index(v_1 = a) < index(v_2 = c)$ is true. The ite operation places then$(f) = b$ on the solid edge of the root of the new ZDD and performs else$(f) + g \pmod 2$, as shown in Fig. 6. During this*

Algorithm 3 Algorithm for performing $f + g \pmod 2$

1: **procedure** $mod\_2\_sum(f, g)$
2:     **if** $f = 0$ **then**
3:         **return** g
4:     **else if** $g = 0$ **then**
5:         **return** f
6:     **else if** $f = g$ **then**
7:         **return** 0
8:     **else**
9:         $v_1 = top\_var(f); v_2 = top\_var(g);$
10:         **if** $index(v_1) < index(v_2)$ **then**
11:             **return** $ite(v_1, then(f), else(f) + g)$
12:         **else if** $index(v_1) > index(v_2)$ **then**
13:             **return** $ite(v_2, then(g), else(g) + f)$
14:         **else**
15:             **return** $ite(v_1, then(f) + then(g), else(f) + else(g))$

---

recursive call, the last condition is true (as $index(v_1 = c) = index(v_2 = c) = 2$). This time the ite operation performs two recursive calls $then(f) + then(g)$ and $else(f) + else(g)$, and so on, to finally construct the ZDD for the Boolean polynomial $ab + d$.

A similar recursive algorithm is also implemented for $f \cdot g \pmod 2$ operation where the intermediate partial product terms are added $\pmod 2$ using the Algorithm 3. *E.g.,* $f = ab, g = a + b, f \cdot g = ab + ab = 0$.

B. GBR $z_i \xrightarrow{G}_+ r_i$ under RTTO on ZDDs

Once the ZDDs for the circuit have been built and stored in G, we need to perform the canonical Gröbner basis reduction $z_i \xrightarrow{G}_+ r_i$ for each output bit z_i. The polynomial r_i will be a canonical representation of z_i in terms of primary inputs. Reduction $f \xrightarrow{g} r$ requires one to obtain the leading monomials $lm(f), lm(g)$ of f, g (Eqn. 7).

1) Finding the leading monomials under RTTO on ZDDs: Recall that RTTO imposes a *lex* term order on the polynomials using the variable order $x_1 > \cdots > x_n$. Moreover, the same variable order $x_1 > \cdots > x_n$ is imposed on the ZDDs. Traversing the then-edges from the root node of a ZDD to terminal **1** delivers the leading monomial of that polynomial under RTTO.

2) Classical reduction with ZDDs: Cancel one monomial in every step: The algorithm for conventional reduction procedure using ZDDs is shown in Algorithm 4. The input parameters are the ZDD of the output bit z_i of the circuit and *poly_list* – a list containing the ZDDs for the set of polynomials $F = \{f_1, \ldots, f_s\}$ corresponding to the gates of the circuit. The algorithm is based on the same principles as the classical division procedure (Algorithm 1).

The procedure $leading\_term(g)$ returns the leading term of the ZDD representation of polynomial g. If g divides f, then the procedure $ZDD\_Divide(f, g)$ (performs cube division) returns the quotient of the division, else it returns zero. Line 8 computes $z_i = z_i + \frac{lt(z_i)}{lt(g)} \cdot g$. The polynomial z_i is completely reduced *w.r.t.* the polynomial g in the while loop.

3) Improved Reduction: Cancel multiple monomials in 1 step: Next, we will show how z_i can be reduced by a

---

Algorithm 4 Reduction: Cancel 1 monomial every iteration

1: **procedure** $single\_mon\_red(z_i, poly\_list)$
2:     **for** each $g \in poly\_list$ **do**
3:         $lead\_g = leading\_term(g)$
4:         $lead\_z_i = leading\_term(z_i)$
5:         $quotient = ZDD\_Divide(lead\_z_i, lead\_g)$
6:         **while** $quotient \neq zero$ **do**
7:             $prod = quotient \cdot g$
8:             $z_i = z_i + prod$
9:             $lead\_z_i = leading\_term(z_i)$
10:             $quotient = ZDD\_Divide(lead\_z_i, lead\_g)$
11:     **return** z_i

---

polynomial g in one step. In the example of Figs. 4 and 5, the primary output z is reduced by f_1 to get r_1. The next step is to reduce r_1 by f_2 to get r_2. To demonstrate our approach we will show how the reduction of r_1 by f_2 can be achieved in one step. There are two monomials in r_1 that contain y, namely yd, y. Both can be canceled by $lt(f_2) = y$ in one step, eliminating the need of the while loop in Algorithm 4.

The polynomial $r_1 = yd + y + d$ can be written as $y \cdot (d + 1) + d$. If we perform 1-step reduction of r_1 by f_2 we get the *quotient* $d + 1$. This quotient is visible as the polynomial represented by the *then*-node of r_1 (Fig. 7). So the reduction can be performed by multiplying $d + 1$ with f_2 and subtracting (adding) this product to $r_1 \pmod 2$.

$$r_1 \xrightarrow{f_2}_+ r_2 \tag{11}$$
$$= (yd + y + d) + (d + 1) \cdot (y + xc + x + c) \pmod 2 \tag{12}$$
$$= 2 \cdot (yd + y) + d + (d + 1) \cdot (xc + x + c) \pmod 2 \tag{13}$$
$$= \underbrace{d}_{else(r_1)} + \underbrace{(d + 1)}_{then(r_1)} \cdot \underbrace{(xc + x + c)}_{else(f_2)} \pmod 2 \tag{14}$$

As shown in Fig. 7, $else(r_1) = d, then(r_1) = d + 1$ and $else(f_2) = xc + x + c$. Moreover, $2 \cdot (yd + y) = 0 \pmod 2$. Therefore, in order to reduce number of operations incurred in the division process, we directly use the last step as a formula for reduction:

$$r_1 \xrightarrow{f_2}_+ r_2 = else(r_1) + then(r_1) \cdot else(f_2) \tag{15}$$

More generally, to perform the division $r_i \xrightarrow{f_j}_+ r_j$, we need to ensure that if f_j divides r_i, then the variable (index) associated with the top-most nodes of their respective ZDDs are the same. This can be ensured by populating $poly\_list = \{f_1, \ldots, f_s\}$ in a certain order.

Due to RTTO, each polynomial $f_j \in F$ is of the form $f_j = x_j + tail(f_j)$, where each variable in $tail(f_j)$ is ordered as being less than x_j (Prop. III.2). Due to this order, the ZDD representation of the polynomials is of the type $f_1 = x_1 + else(f_1), \ldots, f_s = x_s + else(f_s)$, where the variables $x_1, \ldots, x_s$ are the top-most nodes in their respective ZDDs with variable order $x_1 > \cdots > x_s > \cdots > x_n$. Note that the variables $\{x_{s+1}, \ldots, x_n\}$ are the primary inputs, and they are not the output of any logic gate. *We ensure that primary inputs appear last in RTTO.* Then we store elements in *poly_list* according to the order $f_1 > f_2 > \cdots > f_s$; i.e. $poly\_list[1] = f_1, \ldots, poly\_list[s] = f_s$.

Fig. 6: $f + g$ (mod 2) using ZDDs

Fig. 7: ZDDs for polynomial r_1 and f_2.

Considering Algorithm 4, suppose that we first reduce z_i by $poly\_list[1] = f_1$, which results in variable x_1 being replaced by $tail(f_1)$ in z_i. The variable x_1 cannot appear again in z_i at any further reduction step as x_1 is not in the support of polynomials $poly\_list[2], \ldots, poly\_list[s]$. Further in the reduction process, let us assume that we have reduced z_i by $f_1, f_2, \ldots, f_{j-1}$. At this stage the intermediate remainder z_i will not contain any of the variables $x_1, x_2, \ldots, x_{j-1}$ as they have been canceled by the leading terms of $f_1, f_2, \ldots, f_{j-1}$. The next variable in RTTO is x_j. To cancel terms in z_i containing variable x_j, we need to search for the divisor polynomial $f_j = x_j + tail(f_j)$. There can be two possibilities for z_i:

1) If terms in z_i contain x_j, then the top variable of the ZDD of z_i will be x_j. In that case, as the top variable of the ZDD of f_j is also x_j, f_j can divide z_i.

2) If z_i does not contain terms in x_j, then the top variable of the ZDD of z_i will not be x_j. In that case, f_j (with top variable x_j) cannot divide z_i.

Therefore, the divisibility of z_i by the current entry in $poly\_list[j] = f_j$ can be checked just by comparing the indices of the top nodes of the ZDDs of z_i and f_j. If the indices are equal, then f_j divides z_i. Otherwise, f_j does not divide z_i. This allows to replace the cube division (Line 5, Algorithm 4) by an equality check of top indices of ZDDs.

Therefore, unlike in Algorithm 4, where we need to obtain the leading monomials and compute the quotient $lead\_z_i/lead\_g$, now we only need to determine if $lead\_g$ can divide z_i at all (in which case the quotient is $then(z_i)$). This can be accomplished by just comparing the indices of top-most nodes of z_i and g.

It can now be shown that Eqn. (15) holds throughout the GB-reduction under RTTO. To perform the operation $r_i \xrightarrow{f_j}_+ r_j$, we consider their ZDDs constructed using RTTO. Based on the above discussion, the top variables of the ZDDs of r_i

and f_j are the same and equal to x_j. Also, let q denote the quotient of division of r_i by f_j. The operation $r_i \xrightarrow{f_j}_+ r_j$, is carried out as follows,

$$r_j = r_i + q \cdot f_j$$

As x_j is the top variable of r_i and f_j, they can be written as $r_i = x_j \cdot then(r_i) + else(r_i)$ and $f_j = x_j + else(f_j)$, respectively. Also notice that $q = then(r_i)$, because f_j (with leading term x_j) can only divide the $x_j \cdot then(r_i)$ component of r_i.

$$\begin{aligned}
r_j &= x_j \cdot then(r_i) + else(r_i) + q \cdot (x_j + else(f_j)) \\
&= x_j \cdot then(r_i) + else(r_i) + then(r_i) \cdot (x_j + else(f_j)) \\
&= 2 \cdot x_j \cdot then(r_i) + else(r_i) + then(r_i) \cdot else(f_j) \\
&= else(r_i) + then(r_i) \cdot else(f_j)
\end{aligned}$$

The last expression is the same as in Eqn. (15). So the reduction process effectively involves just two operations, a modulo 2 sum and a product. *This has the effect of canceling all the terms in r_i that can be canceled by $lt(f_j)$, implicitly canceling multiple monomials in one step.* This is made possible only due to the properties that RTTO imposes on the structure of ZDDs.

The algorithm for implicit cancellation of multiple monomials in GB-reduction is shown in Algorithm 5, where the notations, z_i and $poly\_list$, are the same as in Algorithm 4. This algorithm significantly reduces the number of iterations, which now exactly equals the size of $poly\_list$. For the example of Fig. 4, the number of iterations is 3 using Algorithm 5, whereas 7 iterations are required using Algorithm 4.

Algorithm 5 Reduction under RTTO: Cancel multiple monomials

1: **procedure** $multi\_mon\_red(z_i, poly\_list)$
2:     **for** each $g \in poly\_list$ **do**
3:         **if** $index(g) == index(z_i)$ **then**
4:             $z_i = else(z_i) + then(z_i) \cdot else(g)$
5:     **return** z_i

We have implemented the above GBR procedures using the CUDD package [17]. The circuit under verification is analyzed, RTTO based variable order is imposed on the ZDDs, and the Boolean polynomials of the circuit are represented as unate cube sets on ZDDs. The polynomials corresponding to the gates of circuit, $G = \{f_1, \ldots, f_s\}$, are inserted in $poly\_list$ according to the variable order $x_1 > \cdots > x_i > \cdots > x_n$, where $f_i = x_i + else(f_i)$. To perform GBR $z_i \xrightarrow{G}_+ r_i$, Algorithm 5 is invoked and r_i used for equivalence checking.

VI. Experimental Results

This section presents the results of using our implementation (Algorithm 5) for formal verification and equivalence checking of circuits used in cryptography. We compare our results against: i) F4-style reduction [6] which models the reduction as Gaussian elimination on a coefficient matrix; ii) Parallelized approach for performing reductions on Galois field multipliers [14]; and iii) PolyBori's [20] reduction procedure with ZDDs as the underlying data structure. For all tools and experiments, RTTO $>$ is used for constraint representation. The experiments are performed on a 3.5GHz Intel Core^{TM} i7-4770K Quad-Core CPU with 32 GB of RAM. The datapath sizes k are selected according to cryptography standards recommended by U.S. National Institute of Standards and Technology (NIST).

A. Mastrovito Multipliers

Modular multiplication is an important computation used in cryptography. A Mastrovito multiplier architecture can be employed for performing this computation over the finite field of 2^k elements, i.e. $\mathbb{F}_{2^k}$. Mastrovito multipliers compute $Z = A \times B \pmod{P(x)}$ where $P(x)$ is a given primitive polynomial for the datapath size k. Here $A = \{a_0, a_1, \ldots, a_{k-1}\}$, $B = \{b_0, b_1, \ldots, b_{k-1}\}$ are the two data operands, and $Z = \{z_0, z_1, \ldots, z_{k-1}\}$ is the output. The product $A \times B$ is first computed using an array multiplier architecture, and then the result is reduced modulo $P(x)$.

Table I provides the results for the reductions $z_i \xrightarrow{G}_+ r_i$ for Mastrovito multipliers for each output bit $z_i, 0 \le i \le k-1$. The benchmarks are taken from [5] and optimized using ABC [28] with the commands *resyn*2 and *dch* as mentioned in [14]. Algorithm 5 reduces each output bit independently of other bits. Therefore, we have presented the results obtained by running our reduction algorithm both sequentially and in parallel for each output bit. Similarly, the results for implementation in PolyBori are also presented for both cases. The implementation presented in [14] is already parallelized. We parallelized PolyBori and our implementation by creating individual processes for each output bit z_i that has its own set of variables and gate polynomials, *poly_list* (Algorithm 5). The maximum number of parallel processes is decided upon the memory usage of each process (*i.e.* reducing one bit) for our implementation and the total available memory. The larger benchmarks are run with fewer parallel processes as they consume more memory.

In the table, the column #T represents the number of parallel processes. (S) and (P) refer to the cases when the experiments run sequentially and in parallel for the output bits z_i, respectively.

The 571-bit multiplier could not be synthesized and mapped with the given memory due to its large size. Therefore, we have provided results for a structured (but unoptimized) 571-bit multiplier benchmark. Our implementation outperforms the explicit approaches of [6] and [14] for Mastrovito multipliers and also PolyBori. For the 571-bit multiplier, the implementation of [6] does not finish for the given time period of 30 hours and the PolyBori implementation crashes (CR). The maximum

TABLE I: Mastrovito Multipliers (Time in seconds); k = Datapath Size, #Gates = No. of gates, #T = No. of threads, Time-Out = 30 hrs, (P): Parallel Execution, (S): Sequential Execution, K = 10^3, M = 10^6, PB: PolyBori, ZR: Algorithm 5

k	#Gates	F4 [6]	#T	[14](P)	PB		ZR	
					(P)	(S)	(P)	(S)
64	11.5K	1.3	20	3.70	3.60	2.21	0.73	**0.27**
128	46K	9.89	20	27.54	23.99	16.76	5.08	**1.63**
163	73.5K	32.61	20	55.96	48.67	33.72	11.41	**3.11**
233	122K	86.30	20	127.61	112.96	77.23	21.77	**3.63**
283	193K	274.68	20	253.05	227.77	157.45	49.89	**11.41**
409	386K	2,528.5	10	716.80	659.64	426.92	163.52	**17.68**
571*	1.6M	TO	3	5,331	CR	CR	2,126.7	**566.4**

memory consumption by ZDDs for Mastrovito multipliers when GBR is performed sequentially varies from 131.6 MB for 64-bit operands to 8.1 GB for 571-bit operands.

An interesting point to note in Table I is that our implementation takes less time when run sequentially. There is a certain overhead involved when we declare variables and build ZDDs for each gate of the circuit. In the case of Mastrovito multiplier benchmarks, this overhead is substantially greater than the reduction time for each output bit. Therefore, when the output bits are reduced in parallel for these benchmarks, this overhead increases the overall run time.

B. Montgomery Multipliers

Exponentiation operations are often required in cryptosystems. For such applications, Montgomery architectures [29] [30] [31] are considered more efficient than Mastrovito multipliers as they do not require explicit reduction modulo P after each step. Fig. 8 shows the structure of a Montgomery multiplier. Each MR block computes $A \cdot B \cdot R^{-1}$, where R is selected as a power of a base (α^k) and R^{-1} is the multiplicative inverse of R in $\mathbb{F}_{2^k}$. As this operation cannot compute $A \cdot B$ directly, we need to pre-compute $A \cdot R$ and $B \cdot R$ as shown in the Fig. 8. We denote the leftmost two blocks as Block A (upper) and B (lower), the middle block as Block C and the output block as Block D.

Fig. 8: Montgomery multiplication.

Table II provides the results for GBR on flattened (bit-blasted) and optimized Montgomery multipliers for the sequential and parallel executions of Algorithm 5. The maximum memory consumption for sequential execution varies from 66.7 MB for 64-bit operands to 11.3 GB for 571-bit operands. Our approach outperforms conventional explicit approaches except for the case of 283 bit multiplier.

Table III presents the statistics for Montgomery multipliers where the hierarchy of Fig. 8 for the blocks A, B, C, and D is made available. The experiment first reduces the outputs of each individual block modulo the gates of that block, and

TABLE II: Montgomery Multipliers (Time in seconds); k = Datapath Size, #Gates = No. of gates, #T = No. of threads, Time-Out = 30 hrs, (P): Parallel Execution, (S): Sequential Execution, K = 10^3, M = 10^6, PB: PolyBori, ZR: Algorithm 5

k	#Gates	F4 [6]	#T	[14](P)	PB		ZR	
					(P)	(S)	(P)	(S)
64	9.5K	16.29	20	10.69	6.27	9.22	**3.75**	8.37
128	35K	621.90	20	36.19	28.93	34.59	**13.76**	24.73
163	56.5K	2,608.4	20	204.94	167.73	335.2	**141.68**	321.60
233	111K	385.92	20	132.51	119.77	99.36	42.16	**31.88**
283	165K	5,344	20	**704.13**	1,194.2	2,078	1,065.3	2,113
409	340K	7,104	10	697.91	737.23	722.1	303.91	**299.92**
571*	1.97M	TO	3	TO	CR	CR	**43,813**	99,042

then reduces the primary outputs modulo these four sets of remainders (ZDDs), thus exploiting the hierarchy of these circuits. Table III shows the time for reduction of each block and the time for reducing the primary outputs across the four blocks. The time for reducing the primary outputs across the hierarchical blocks in case of the F4 implementation is <1 second, and is not explicitly mentioned in the table. The row labeled *Total* presents the sum of the computation time of reduction across these levels, and the maximum of the time to reduce each MR block (as the reductions for the four blocks are independent of each other and are parallelized). These results again demonstrate the efficiency of our approach against explicit approaches.

TABLE III: Montgomery Blocks (Time in seconds); k = Datapath Size, #Gates = No. of gates, Time-Out = 30 hrs, Red. = time for reduction, Coll. = time to reduce across the 4 levels. K = 10^3, M = 10^6, PB: PolyBori, ZR: Algorithm 5

k	#Gates	Block	F4 [6]	PB		ZR	
				Red.	Coll.	Red.	Coll.
163	33K	Block A	25	12		1	
	33K	Block B	25	12	16	1	18
	85K	Block C	73	18		7	
	32K	Block D	24	12		1	
		Total	73	34		**25**	
233	55K	Block A	142	32		0.14	
	55K	Block B	141	33	5	0.14	4
	163K	Block C	408	34		2.1	
	54K	Block D	140	32		0.13	
		Total	408	39		**6.1**	
283	82K	Block A	330	79		24	
	82K	Block B	329	78	26	23	90
	241K	Block C	883	173		118	
	81K	Block D	321	80		23	
		Total	883	**199**		208	
409	168K	Block A	1,322	177		0.57	
	168K	Block B	1,335	175	28	0.57	29
	502K	Block C	4,471	192		14	
	168K	Block D	1,338	176		0.56	
		Total	4,471	220		**43**	
571	330K	Block A	5,371	769		321	
	330K	Block B	5,421	747	1,341	332	1,412
	980K	Block C	37,804	3,605		3026	
	328K	Block D	5,539	751		338	
		Total	37,804	4,946		**4,438**	

Equivalence Checking: As a result of the GBR $z_i \xrightarrow{G}_+ r_i$, the function implemented by each output bit z_i of the circuit is represented as a reduced, canonical, Boolean polynomial in terms of the primary inputs, and by using a ZDD. Thus

the equivalence of such vastly different arithmetic circuit implementations (Mastrovito vs Montgomery) can be verified by testing for the equality (isomorphism) of the corresponding ZDD graphs.

C. Point Addition over Elliptic Curves

Point addition is an important operation required for the task of encryption, decryption and authentication in Elliptic Curve Cryptography (ECC). Modern approaches represent the points in projective coordinate systems, *e.g.*, the López-Dahab (LD) projective coordinate [32], due to which the point addition operation can be implemented as polynomials in the field.

Example VI.1. *Consider point addition in López-Dahab (LD) projective coordinate. Given an elliptic curve:* $Y^2 + XYZ = X^3Z + aX^2Z^2 + bZ^4$ *over* $\mathbb{F}_{2^k}$, *where* X, Y, Z *are k-bit vectors that are elements in* $\mathbb{F}_{2^k}$ *and similarly, a, b are constants from the field. We represent point addition over the elliptic curve as* $(X_3, Y_3, Z_3) = (X_1, Y_1, Z_1) + (X_2, Y_2, 1)$. *Then X_3, Y_3, Z_3 can be computed as follows:*

$$A = Y_2 \cdot Z_1^2 + Y_1 \qquad B = X_2 \cdot Z_1 + X_1$$
$$C = Z_1 \cdot B \qquad D = B^2 \cdot (C + aZ_1^2)$$
$$Z_3 = C^2 \qquad E = A \cdot C$$
$$X_3 = A^2 + D + E \qquad F = X_3 + X_2 \cdot Z_3$$
$$G = X_3 + Y_2 \cdot Z_3 \qquad Y_3 = E \cdot F + Z_3 \cdot G$$

Each of the polynomials in the above design are implemented as a (gate-level) logic block and are interconnected to obtain final outputs X_3, Y_3 and Z_3.

TABLE IV: Point Addition Circuits (Time in seconds); k = Datapath Size, #Gates = No. of gates, Time-Out = 30 hrs, K = 10^3, M = 10^6, PB: PolyBori, ZR: Algorithm 5

k	#Gates	F4 [6]	PB	ZR
64	15.3K	1.78	3.32	**0.72**
128	64K	40.55	27.41	**6.03**
163	104K	130.24	57.57	**13.13**
233	139K	335.60	106.85	**19.62**
283	281K	1,787.96	273.53	**64.48**
409	423K	5,077.50	578.15	**115.20**
571	1.14M	48,162.29	CR	**725.95**

The word-level abstraction approach in [6] presents the results for extracting the above representation for each of $A, B, \ldots, X_3, Y_3, Z_3$ blocks. It first performs a bit-level reduction for every output of each block (GBR $z_i \xrightarrow{G}_+ r_i$), and then a bit-to-word substitution to derive an input-output word-level representation for the circuit. Table IV shows the comparison of the time required for bit-level reduction of outputs d_i of the block $D = B^2 \cdot (C + aZ_1^2)$ as done in [6] against our implementation. (Bit-level reductions for other blocks take much less time than that for block D.) This result demonstrates that our bit-level GBR implementation is in many cases orders of magnitude faster than the F4-style reduction of [6]. Therefore our approach can replace the F4-style bit-level GBR of [6] and improve the overall process of word-level abstraction of datapath designs.

D. Equivalence Checking of Sequential Galois Field Multipliers

The designs discussed so far are combinational implementations of polynomial computations of finite filed circuits. These designs use the standard basis representation $\{1, \alpha, \alpha^2, \ldots, \alpha^{k-1}\}$ to model a k-bit data-word Z in terms of its constituent bits as $Z = z_0 + z_1\alpha + z_2\alpha^2 \cdots + z_{k-1}\alpha^{k-1}$, with α being the primitive element for that field $\mathbb{F}_{2^k}$.

There exists sequential multipliers where k-bit inputs are loaded into k-bit registers, and the k-bit result is available after k clock-cycle execution of the machine. These multipliers use a *normal basis* $\{\beta, \beta^2, \beta^4, \ldots, \beta^{2^{k-1}}\}$ to represent a k-bit data-word S in terms of its constituent bits as $S = s_0\beta + s_1\beta^2 + s_2\beta^4 \cdots s_{k-1}\beta^{2^{k-1}}$, with β being the normal element. The relation between α and β can be used to represent the bits z_i and s_j in terms of each other.

We perform equivalence checking between two different architectures of *sequential multipliers with parallel output (SMPO)*, the Agnew-SMPO (AG-SMPO) by G.B. Agnew [33] and the RH-SMPO by Reyhani-Masoleh and Hasan [34]. In order to perform equivalence checking, the circuits are unrolled over k time frames, and the GBR $s_i \xrightarrow{G}_+ r_i$ is performed to obtain a canonical r_i (G is the set of polynomials for the unrolled circuit under RTTO). The ZDDs for respective r_i's (for AG-SMPO and RH-SMPO) are compared to perform equivalence check.

Tables V and VI present the run-time of our implementation for performing these reductions on RH-SMPO and AG-SMPO architectures respectively, when compared with the approach presented in [6] and PolyBori. The results show that our implementation is about an order of magnitude faster than PolyBori and multiple orders of magnitude faster than the explicit approach of [6].

TABLE V: RH-SMPO Multipliers (Time in seconds); k = Datapath Size, #Gates = No. of gates, Time-Out = 30 hrs, K = 10^3, PB: PolyBori, ZR: Algorithm 5

$k =$	65	81	89	131	173	233	281	410
#Gates	13.6K	21.4K	25.9K	55.9K	96.5K	177K	258K	546K
F4[6]	9.02	26.65	42.46	294.7	874.3	3,404	7,328	23,610
PB	3.65	6.07	7.42	28.22	47.16	116.63	199.32	637.69
ZR	0.42	0.80	1.01	3.03	3.53	8.12	13.27	52.09

TABLE VI: AG-SMPO Multipliers (Time in seconds); k = Datapath Size, #Gates = No. of gates, Time-Out = 30 hrs, K = 10^3, PB: PolyBori, ZR: Algorithm 5

$k =$	65	81	89	131	173	233	281	410
#Gates	12.5K	19.5K	23.6K	51.2K	89.4K	162K	236K	503K
F4[6]	8.34	20.46	33.2	221.4	754.1	2,655	5,569	21,938
PB	3.11	6.82	9.21	20.15	44.37	107.12	187.77	578.61
ZR	0.44	0.77	0.91	2.51	3.39	7.8	12.63	43.78

E. Limitations of our approach: Integer arithmetic circuits

The GBR approach presented in the previous section, although applicable to integer arithmetic circuits, is not computationally feasible for their verification. We evaluated our technique on integer arithmetic multiplier circuits, which showed an exponential increase in verification time *w.r.t.* the circuit size. This is because our approach reduces each primary output bit independently, and does not consider the logic sharing among different primary outputs. As a result, a large number of monomials are generated during reduction that are common to multiple primary outputs. This effect of logic sharing cannot be exploited in GBR when each primary output bit is reduced independently. For example, the ZDD-based bit-level GBR for a 7x7 integer multiplier reveals that when reducing the z_{13} bit (MSB) and the z_{12} bit independently, the maximum number of monomials encountered (*i.e.* during $z_{13} \xrightarrow{G}_+ r_{13}$ and $z_{12} \xrightarrow{G}_+ r_{12}$) are 429,889 and 897,955, respectively. However, the modulo-2 sum (XOR) of these ZDDs contains only 789,604 monomials (during the modulo-2 sum common monomials cancel out) as opposed to 1,327,844 (= 429,889 + 897,955). A word-level reduction approach (such as those of [9] and [11]) may cancel such common non-linear terms early in the reduction process and avoid intermediate blow-up in the number of monomials. This is shown below.


Fig. 9: Integer multiplier circuit

Consider the integer multiplier circuit given in Fig. 9. A word-level approach would model the output word as $Z = z_0 + 2z_1 + 4z_2 + 8z_3$, and perform the reduction $Z \xrightarrow{G}_+ R$ across the reverse topological levels (RTTO) depicted in the figure. The polynomials of the circuit are:

$$z_0 = x_4$$
$$z_1 = x_2 + x_3 - 2x_2x_3$$
$$z_2 = x_1 + x_5 - 2x_1x_5 = x_1 + x_2x_3 - 2x_1x_2x_3$$
$$z_3 = x_1x_5 = x_1x_2x_3$$

Here $+$ denotes addition over integers. The reduction of the word-level expression $8z_3 + 4z_2 + 2z_1 + z_0$ in Z cancels out the common nonlinear monomials to control intermediate monomial explosion:

$$
\begin{aligned}
Z &= 8z_3 + 4z_2 + 2z_1 + z_0 \\
&= 8x_1x_2x_3 + (4x_1 + \underbrace{4x_2x_3} - \underline{8x_1x_2x_3}) \\
&\quad + (2x_2 + 2x_3 - \underbrace{4x_2x_3}) + x_4 \\
&= 4x_1 + 2x_2 + 2x_3 + x_4
\end{aligned}
$$

A purely bit-level GBR approach is thus not suitable for integer arithmetic circuits. However, this is not a limitation of

our algorithms and implementations, but rather an issue of the capability of bit-level versus word-level models. Integrating the implicit data structure with a word-level representation may yield better results for such applications.

VII. CONCLUSION

This paper has presented an approach for formal verification of datapath circuits by deriving a canonical polynomial representation for each output bit z_i of a circuit in terms of the primary inputs using Gröbner basis reduction. The gates of the circuit C are modeled as a set of polynomials G over $\mathbb{F}_2$ where the variables are the nets of the circuit. An order on the variables is derived from the topology of the circuit, and a *lex* term order (RTTO) is imposed on the polynomials. RTTO renders the set G itself a Gröbner basis. The reduction $z_i \xrightarrow{G}_+ r_i$ results in a canonical remainder r_i for each output z_i.

The polynomials in the set G are Boolean polynomials that can be construed as unate cube sets. The unate cube set algebra prowess of ZDDs is exploited to represent the polynomials implicitly. We show that RTTO imposes a special structure on the ZDDs, where subexpressions for leading monomials and quotients of the division are readily visible as subgraphs in the ZDD. We take further advantage of this data structure to improve the classical Gröbner basis reduction method that relies on canceling only 1 monomial in every iteration of division. Our approach cancels multiple monomials in each step of division and generates fewer terms, thus speeding up the reduction. We have performed experiments with various finite field circuits used in cryptography. Our approach achieves significant improvement over recent approaches: the F4-style reduction, a parallelized approach for reduction, and PolyBori.

REFERENCES

[1] R. E. Bryant, "Graph Based Algorithms for Boolean Function Manipulation," *IEEE Trans. on Computers*, vol. C-35, pp. 677–691, August 1986.

[2] A. Mishchenko, S. Chatterjee, R. Brayton, and N. Een, "Improvements to Combinational Equivalence Checking," in *Proc. Intl. Conf. on CAD (ICCAD)*, 2006, pp. 836–843.

[3] O. Wienand, M. Wedler, D. Stoffel, W. Kunz, and G. Gruel, "An Algebraic Approach to Proving Data Correctness in Arithmetic Datapaths," in *Computer Aided Verification Conf.*, 2008, pp. 473–486.

[4] E. Pavlenko, M. Wedler, D. Stoffel, W. Kunz, A. Dreyer, F. Seelisch, and G.-M. Greuel, "STABLE: A New QBF-BV SMT Solver for Hard Verification Problems Combining Boolean Reasoning with Computer Algebra," in *Proc. Des. Auto. Test in Europe Conf.*, 2011, pp. 155–160.

[5] J. Lv, P. Kalla, and F. Enescu, "Efficient Gröbner Basis Reductions for Formal Verification of Galois Field Arithmetic Circuits," in *IEEE Trans. on CAD*, vol. 32, no. 9, 2013, pp. 1409–1420.

[6] T. Pruss, P. Kalla, and F. Enescu, "Efficient Symbolic Computation for Word-Level Abstraction from Combinational Circuits for Verification over Finite Fields," *IEEE Trans. on CAD*, vol. 35, no. 7, pp. 1206–1218, July 2016.

[7] W. W. Adams and P. Loustaunau, *An Introduction to Grobner Bases*. American Mathematical Society, 1994.

[8] X. Sun, P. Kalla, and F. Enescu, "Word-level Traversal of Finite State Machines using Algebraic Geometry," in *Proc. High-Level Design Validation and Test*, 2016, pp. 142–149.

[9] M. Ciesielski, C. Yu, W. Brown, D. Liu, and A. Rossi, "Verification of Gate-level Arithmetic Circuits by Function Extraction," in *Proc. Design Automation Conference (DAC)*, 2015, pp. 1–6.

[10] F. Farahmandi and B. Alizadeh, "Groebner Basis Based Formal Verification of Large Arithmetic Circuits Using Gaussian Elimination and Cone-based Polynomial Extraction," *Microprocessors and Microsystems*, vol. 39, no. 2, pp. 83–96, 2015.

[11] A. Sayed-Ahmed, D. Große, U. Kühne, M. Soeken, and R. Drechsler, "Formal Verification of Integer Multipliers by Combining Gröbner Basis with Logic Reduction," in *Proc. Des. Auto. Test in Europe Conf. (DATE)*, 2016, pp. 1048–1053.

[12] A. S. Ahmed, D. Große, M. Soeken, and R. Drechsler, "Equivalence Checking Using Gröbner Bases," in *Formal Methods in Computer-Aided Design (FMCAD)*, 2016, pp. 169–176.

[13] D. Ritirc, A. Biere, and M. Kauers, "Column-Wise Verification of Multipliers Using Computer Algebra," in *Formal Methods in Computer-Aided Design (FMCAD)*, 2017, pp. 23–30.

[14] C. Yu and M. Ciesielski, "Efficient Parallel Verification of Galois Field Multipliers," in *Proc. Asia and South Pacific Des. Auto. Conf.*, 2017, pp. 238–243.

[15] S. Minato, "Zero-Suppressed BDDs for Set Manipulation in Combinatorial Problems," in *Design Automation Conference (DAC)*, 1993, pp. 272–277.

[16] W. Decker, G.-M. Greuel, G. Pfister, and H. Schönemann, "SINGULAR 3-1-6 — A computer algebra system for polynomial computations," http://www.singular.uni-kl.de, 2012.

[17] F. Somenzi, "CUDD: CU Decision Diagram Package Release 3.0.0," http://vlsi.colorado.edu/~fabio, 2015, University of Colorado at Boulder.

[18] N. Shekhar, P. Kalla, F. Enescu, and S. Gopalakrishnan, "Equivalence Verification of Polynomial Datapaths with Fixed-Size Bit-Vectors using Finite Ring Algebra," in *Intl. Conf. on Computer-Aided Design, ICCAD*, 2005.

[19] O. M. Hansen and J.-F. Michon, "Boolean Gröbner basis," in *Proc. Boolean Functions Cryptography & Applications*, 2006, pp. 185–201.

[20] M. Brickenstein and A. Dreyer, "Polybori: A Framework for Gröbner Basis Computations with Boolean Polynomials," *Journal of Symbolic Computation*, vol. 44, no. 9, pp. 1326–1345, September 2009.

[21] M. Clegg, J. Edmonds, and R. Impagliazzo, "Using the Gröbner Basis Algorithm to Find Proofs of Unsatisfiability," in *ACM Symposium on Theory of Computing*, 1996, pp. 174–183.

[22] G. Avrunin, "Symbolic Model Checking using Algebraic Geometry," in *Computer Aided Verification Conf.*, 1996, pp. 26–37.

[23] M. Y. Vardi and Q. Tran, "Groebner Bases Computation in Boolean Rings for Symbolic Model Checking," in *Proc. IASTED Intl. Conf. on Modelling and Simulation*, 2007, pp. 440–445.

[24] B. Buchberger, "Ein Algorithmus zum Auffinden der Basiselemente des Restklassenringes nach einem nulldimensionalen Polynomideal," Ph.D. dissertation, Philosophiesche Fakultat an der Leopold-Franzens-Universitat, Austria, 1965.

[25] S. Gao, "Counting Zeros over Finite Fields with Gröbner Bases," Master's thesis, Carnegie Mellon University, 2009.

[26] B. Buchberger, "A Criterion for Detecting Unnecessary Reductions in the Construction of a Gröbner Bases," in *Symbolic and Algebraic Computation*, 1979, pp. 3–21.

[27] S. Minato, "Calculation of Unate Cube Set Algebra using Zero-Suppressed BDDs," in *Proc. Design Automation Conference (DAC)*, 1994, pp. 420–424.

[28] Berkeley Logic Synthesis and Verification Group, "ABC: A system for sequential synthesis and verification," www.eecs.berkeley.edu/alanmi/abc, 2007.

[29] C. Koc and T. Acar, "Montgomery Multiplication in GF(2^k)," *Designs, Codes and Cryptography*, vol. 14, no. 1, pp. 57–69, Apr. 1998.

[30] H. Wu, "Montgomery Multiplier and Squarer for a Class of Finite Fields," *IEEE Transactions On Computers*, vol. 51, no. 5, May 2002.

[31] M. Knežević, K. Sakiyama, J. Fan, and I. Verbauwhede, "Modular Reduction in GF(2^n) Without Pre-Computational Phase," in *Proceedings of the International Workshop on Arithmetic of Finite Fields*, 2008, pp. 77–87.

[32] J. López and R. Dahab, "Improved Algorithms for Elliptic Curve Arithmetic in GF(2^n)," in *Proc. Selected Areas in Cryptography*, 1999, pp. 201–212.

[33] G. B. Agnew, R. C. Mullin, I. Onyszchuk, and S. A. Vanstone, "An implementation for a fast public-key cryptosystem," *Journal of CRYPTOLOGY*, vol. 3, no. 2, pp. 63–79, 1991.

[34] A. Reyhani-Masoleh and M. A. Hasan, "Low complexity word-level sequential normal basis multipliers," *Computers, IEEE Transactions on*, vol. 54, no. 2, pp. 98–110, 2005.

PLACE PHOTO HERE

Utkarsh Gupta Utkarsh Gupta received the B.Tech. degree in Electronics & Communication Engineering from Tezpur University in India in 2014. He has been a student at the University of Utah since 2014 where he is currently working towards his PhD. degree in Computer Engineering. His research interests include computer algebra based formal hardware verification, graph based verification, and logic synthesis.

PLACE PHOTO HERE

Priyank Kalla Priyank Kalla received the B.E. degree in Electronics Engineering from the Birla Vishvakarma Mahavidyalaya, Saradar Patel University in India in 1993, and M.S. and PhD. degrees from the University of Massachusetts Amherst in 1998 and 2002, respectively. He has been a faculty member in the Electrical and Computer Engineering Department at the University of Utah since 2002, where he is currently an Associate Professor. His research interests are in fundamental CAD techniques for synthesis, optimization, and verification of digital VLSI circuits and systems. Dr. Kalla was the chair of the IEEE High-Level Design Validation and Test Workshop (HLDVT) 2009, and also the chair of the IEEE CAS Technical Committee on Computer-Aided Network Design (CANDE) 2011-2012. He is a recipient of the NSF CAREER award in 2006, and the ACM TODAES 2009 Best Paper award.

PLACE PHOTO HERE

Vikas Rao Vikas Rao received the B.E. degree from JSSATE, Bangalore, India, in 2008, and the M.S. degree from Manipal University, India, in 2010. From 2009 to 2016, he was a Graphics System Validation Engineer with VPG group, Intel Corporation, Bangalore, India. He is currently pursuing his doctoral degree in computer engineering from the University of Utah. His current research interests include computer-aided design techniques for logic synthesis and formal hardware verification.